



Vrije Universiteit Brussel

Faculty of Science and Bio-Engineering Sciences  
Department of Computer Science  
and Applied Computer Science

# Expressing and checking application-specific, user-specified security policies

---

Graduation thesis submitted in partial fulfillment of the requirements for the degree of  
Master of Science in Applied Sciences and Engineering: Computer Science

**Valentijn Spruyt**

---

Promotor: Prof. Dr. Coen De Roover

Advisors: Jens Nicolay  
Quentin Stievenart

JUNE 2016





Vrije Universiteit Brussel

Faculteit Wetenschappen en Bio-  
Ingenieurswetenschappen  
Departement Computerwetenschappen  
en Toegepaste Informatica

# Expressing and checking application-specific, user-specified security policies

Proefschrift ingediend met het oog op het behalen van de graad van  
Master of Science in Applied Sciences and Engineering: Computer Science

**Valentijn Spruyt**

Promotor: Prof. Dr. Coen De Roover  
Begeleiders: Jens Nicolay  
Quentin Stievenart

JUNI 2016



# **Abstract**

# **Acknowledgements**

# Contents

<b>1</b>	<b>Introduction</b>	<b>5</b>
1.1	Context . . . . .	5
1.2	Motivation . . . . .	5
1.3	Objective . . . . .	5
1.4	Overview . . . . .	5
<b>2</b>	<b>Background</b>	<b>6</b>
2.1	Program representations and querying mechanisms . . . . .	6
2.1.1	Exploring and Enforcing Security Guarantees via Program Dependence Graphs . . . . .	6
2.1.2	GATEKEEPER: Mostly Static Enforcement of Security and Reliability Policies for JavaScript Code . . . . .	6
2.1.3	Parametric regular path queries . . . . .	6
2.2	Expressing policies using a domain-specific language . . . . .	6
2.2.1	Fluent Interfaces to a Java-Based Internal Domain-Specific Languages for Graph Generation and Analysis . . . . .	7
2.2.2	A Little Language for Surveys: Constructing an Internal DSL in Ruby . . . . .	8
2.2.3	Dagoba: an in-memory graph database . . . . .	10
2.3	Conclusion . . . . .	10
<b>3</b>	<b>Context</b>	<b>11</b>
3.1	Static analysis . . . . .	11
3.2	Conclusion . . . . .	11
<b>4</b>	<b>The JS-QL query language</b>	<b>12</b>
4.1	The query language . . . . .	12
4.1.1	Motivation . . . . .	12
4.1.2	Syntax and structure . . . . .	12
4.1.3	Defining policies . . . . .	12
4.2	Conclusion . . . . .	12

<b>5</b>	<b>The query engine</b>	<b>13</b>
5.1	Architecture . . . . .	13
5.2	Types of queries . . . . .	13
5.3	Recursion . . . . .	13
5.4	Conclusion . . . . .	13
<b>6</b>	<b>Implementation</b>	<b>14</b>
6.1	Used technologies . . . . .	14
6.2	Design of the query system . . . . .	14
6.3	Conclusion . . . . .	14
<b>7</b>	<b>Evaluation</b>	<b>15</b>
7.1	the GateKeeper language . . . . .	15
7.2	The PidginQL language . . . . .	15
7.2.1	Only CMS administrators can send a message to all CMS users . . . . .	15
7.2.2	A database is opened only after the master password is checked or when creating a new database . . . . .	18
7.2.3	Public outputs do not depend on a users's password, un- less it has been cryptographically hashed . . . . .	18
7.3	The Conscript language . . . . .	18
7.4	Evaluation . . . . .	18
<b>8</b>	<b>Conclusion and future work</b>	<b>19</b>
8.1	Summary . . . . .	19
8.2	Future work . . . . .	19

# **Chapter 1**

## **Introduction**

### **1.1 Context**

### **1.2 Motivation**

### **1.3 Objective**

### **1.4 Overview**

# Chapter 2

## Background

### 2.1 Program representations and querying mechanisms

Programs can be represented in several ways. These representations can then be queried to detect all kinds of information, such as control- and data-flow properties. In this section, three program representation approaches are presented, together some means to query those representations.

#### 2.1.1 Exploring and Enforcing Security Guarantees via Program Dependence Graphs

#### 2.1.2 GATEKEEPER: Mostly Static Enforcement of Security and Reliability Policies for JavaScript Code

#### 2.1.3 Parametric regular path queries

### 2.2 Expressing policies using a domain-specific language

This section describes three internal domain-specific languages (*DSLs*). We present one DSL written in Java, a statically typed language, and two DSLs written in dynamically typed languages, namely Ruby and JavaScript.



### 2.2.1 Fluent Interfaces to a Java-Based Internal Domain-Specific Languages for Graph Generation and Analysis

Many complex systems problems manifest themselves as networks. Reasoning about these networks can be hard to do manually and asks for complex algorithms to perform sometimes even simple calculations. Hawick[5] pleads for the use of some sort of abstraction to perform graph generation and analysis, more specifically the use of an internal domain-specific language. He presents a DSL built using fluent interface techniques and the statically typed Java programming language. Common data structures and repetitive computations often offer an opportunity to abstract over them, as is the case for models based on networks and graphs.

The goal of this graph DSL is to be able to compare individual network sets to detect characteristic signature properties. The approach is powerful because a major set of data structures and operations on graphs can be abstracted into a library framework, which can then be used by domain experts.

The first step in setting up the DSL is to set up the common data structures. For the graph DSL, there are three: `Nodes`, `Arcs` and of course the `Graphs` themselves. The fields of these data structures are divided into several categories, as depicted in figure 2.1. Structural fields hold the main graph structure, whereas auxiliaries just exist to facilitate computations. Convenience fields contain information that might come in handy, but isn't necessarily used for computations. Finally, decorative fields just are there to have some means of presenting the data in a clear, distinguishable way.

```
class Node{
    // Structural:
    List<Arc> inputs  = new Vector<>();
    List<Arc> outputs = new Vector<>();

    // Convenience:
    List<Node> dsts = new Vector<>();
    List<Node> srcs = new Vector<>();

    // Decorative:
    int index      = 0;
    int mark       = 0;
    double weight  = 1.0;
    String label   = "";

    // Computation Auxiliaries:
    int component  = 0;
    int betweenness = 0;
    int count      = 0;
    boolean visited = false;
    boolean blocked = false;
}
```

Figure 2.1: The 'Node' data structure

Since we now have all information to perform most of the complex computations, the fluent interface can be set up. The approach used here implements the (Java) method chaining technique. This enables the cascading of methods by making each method in the fluent interface return the reference to itself, namely the *this* reference.

Most internal DSL's can be used as a standalone language, as seen in figure 2.2, but the paper also gives some examples in which the DSL is used inside the host language, such as the repetitive removal of the most stressed node in the network to investigate network robustness.

```
public static void main( String args[] ){  
  
    Graph g = Graph.New( args )  
    .setLogging( true )  
    .report ()  
    .removeLeaves ()  
    .computeDegrees ()  
    .computeClusteringCoefficient ()  
    .computeAdjacency ()  
    .computeComponents ()  
    .computePaths ()  
    .computeBetweenness ()  
    .computeDistances ()  
    .computeCircuits ()  
    .report ()  
    .write( "composite.graph" )  
    ;  
}
```

Figure 2.2: Example use of the internal DSL

### 2.2.2 A Little Language for Surveys: Constructing an Internal DSL in Ruby

A Little Language for Surveys [2] explores the use of the Ruby programming language to implement an internal domain-specific language. It checks how well the flexible and dynamic nature of the language accomodates for the implementation of a DSL for specifying and executing surveys. Two key features of the Ruby programming languages are exploited because they especially support defining internal DSLs : The flexibility of the syntax[1] and the support for blocks[4]. Figure 2.3 shows how function calls are easily readable, since the braces surrounding the arguments can be omitted and the arguments list can consist of a variable number of arguments (The latter is also supported in JavaScript[6]). It also shows how entire blocks can be attached to method calls. These blocks are passed unevaluated to the called method, enabling *deferred evaluation*.

```

question "What is your gender?" do
  response "Female" { @female = true }
  response "Male" { @female = false }
  action { @male = if @female then false
              else true end }
end

```

Figure 2.3: Ruby method call syntax

A handy feature in programming languages is reflexive metaprogramming. The survey DSL makes use of the following Ruby reflexive metaprogramming facilities:

**obj.instance\_eval(str)** takes a string `str` and executes it as Ruby code in the context of `obj`. This method allows internal DSL code from a string or file to be executed by the Ruby interpreter.

**mod.class\_eval(str)** takes a string `str` and executes it as Ruby code in the context of module `mod`. This enables new methods and classes to be declared dynamically in the running program.

**obj.method\_missing(sym, \*args)** is invoked when there is an attempt to call an undefined method with the name `sym` and argument list `args` on the object `obj`. This enables the object to take appropriate remedial action.

**obj.send(sym, \*args)** calls method `sym` on object `obj` with argument list `args`. In Ruby terminology, this sends a message to the object.

The design of the survey language is fairly simple. A survey consists of a title, some questions, some responses and finally a result. Each of these actions have a corresponding method in the DSL.

The developers of the survey language chose to split up the parsing and interpretation logic, following the *two-pass architecture*. The first-pass layer parses the file (which is read using `instance_eval`) and generates an abstract syntax tree. These parser classes are structured according to the *Object Scoping* pattern[3], using an approach called *sandboxing*. This architecture is depicted in figure 2.4. The `SurveyBuilder` class evaluates the statements it reads from the input file using it's superclass' methods. This evaluation parses the input file and builds the AST, which is stored in the superclass as well. In this way, the object scoping pattern is applied: All calls are directed to a single object (the superclass), and global namespace cluttering is avoided. When creating the AST nodes, all blocks that were passed to the method calls inside the top-level call are stored in the AST

nodes, instead of being evaluated directly. These blocks will only be evaluated in the second-pass phase (hence *deferred evaluation*). This is illustrated in figure 2.3, where the block passed to `question` is evaluated in the first-pass layer, and the blocks passed to `response` and `action` are stored in the AST, ready to be evaluated in the second-pass layer. Note that sandboxing occurs by calling `instance_eval` inside the `SurveyBuilder` object. In this way, harm can only be done *inside* this object.




Figure 2.4: Ruby method call syntax

The actual interpretation of the created AST happens in the second-pass layer. This is a design decision which allows for different interpretation layers to be plugged in/swapped on-the-fly. However, for the survey language a simple *visitor* pattern implementation suffices to process the AST. Every AST node must provide an `accept` method which takes a `SurveyVisitor` as an argument. This is the only condition that has to be met by the interpretation layer. In this concrete example, there could be a `SurveyConsoleVisitor` and a `SurveyGUIVisitor` class, each representing the survey in their own specific way.

### 2.2.3 Dagoba: an in-memory graph database

## 2.3 Conclusion

# **Chapter 3**

## **Context**

### **3.1 Static analysis**

### **3.2 Conclusion**

# **Chapter 4**

## **The JS-QL query language**

### **4.1 The query language**

#### **4.1.1 Motivation**

#### **4.1.2 Syntax and structure**

#### **4.1.3 Defining policies**

### **4.2 Conclusion**

# **Chapter 5**

## **The query engine**

### **5.1 Architecture**

### **5.2 Types of queries**

### **5.3 Recursion**

### **5.4 Conclusion**

# **Chapter 6**

## **Implementation**

**6.1 Used technologies**

**6.2 Design of the query system**

**6.3 Conclusion**



# Chapter 7

## Evaluation

In this chapter we validate and evaluate the expressiveness of the JS-QL query language by expressing some existing security policies, described in other related work, in our own query language. We will then compare these policies in terms of expressiveness and flexibility. The concept and approach for creating a new, domain-specific language for security policies is explained in chapter 4. Chapter 5 discusses the underlying query engine and how it works together with the query language to process the application-specific policies. In chapter 6 we explain how our approach was instantiated.

We start this chapter by expressing 9 security policies distilled from 3 papers in sections 7.2, 7.1 and 7.3 respectively. Every JS-QL policy will be evaluated by comparing how well it matches the policy expressed in the original paper. Finally, in section 7.4, we evaluate the query framework by specifying its advantages, disadvantages and limitations. We will also briefly compare the query languages presented in this chapter in terms of expressiveness, verbosity and conciseness (*LOC*).

### 7.1 the GateKeeper language

### 7.2 The PidginQL language

In this chapter we attempt to express 3 policies originally presented in [7].

#### 7.2.1 Only CMS administrators can send a message to all CMS users

Imagine a situation where not only administrators can send broadcast messages. A regular user with bad intentions could easily take advantage of this situation to

cause harm to the system. A CMS application for instance with a decent size of users could be exploited by sending a message to all users, asking them to reply with their password. When the attacker provides a reason to the victims convincing them to send their password, he could possibly compromise the contents of the victim's account. An example of such a reason could be that the 'administrator' needs to have the password of a user account in order to update the software of that user to the latest version. This behaviour is undesirable, thus we need a policy which prevents regular users from sending such messages.

The policy listed in [7] that addresses this issue can be found in listing 7.1. First, all nodes that are entries of the `addNotice` method are searched for and stored in a variable. `addNotice` is the method that sends messages to all users, and has the same behaviour as the broadcast method in the explanation above. Next, all points in the PDG are found that match a return node of the `isCMSAdmin` method with a return value which is truthy. In order to know if there exists some path in the graph where `addNotice` is called when the return value of `isCMSAdmin` is false, all paths between the nodes in `addNotice` and `isAdmin` are removed from the graph for all paths where `isAdmin` is true. Finally, the intersection of the nodes in this 'unsanitized' graph and the nodes in the `sensitiveOps` argument is taken. When this intersection is not empty, we can assume that there is a violation of the policy in the remainder of the graph. This last part is exactly what the `accessControlled` method does.

---

```

1 let accessControlled(G, checks, sensitiveOps) =
2   G.removeControlDeps(checks) ∩ sensitiveOps is empty
3
4 let addNotice = pgm.entriesOf("addNotice") in
5 let isAdmin   = pgm.returnsOf("isCMSAdmin") in
6               pgm.accessControlled(isAdmin, addNotice)

```

---

Listing 7.1: Policy 1 in PidginQL

When attempting to write a similar query in JS-QL, we need to define the problem in terms of control flow: "There must be no path between the returns of the `isCMSAdmin`, when the return value is false, and a call of the `addNotice` method." We must note that with abstract interpretation, it is not trivial to specify whether a value is truthy or falsy. When looking at a conditional (like an `if`-statement), we can determine whether the true- or false-branch has been taken by comparing the first node of the branches with the alternate/consequent of the conditional. However, for values with the value of  $\{Bool\}$ , we cannot decide on which branch is the true-branch and which one is the false-branch. We can solve this in two ways: We can assume that the condition in the conditional is a direct call to `isCMSAdmin`, which enables us to find the false-branch. From there on we can search for all calls to `addNotice` to find violations. The JS-QL

policy for this case is defined in listing 7.2. We skip all states until we reach the beginning of a false branch of a conditional. We bind the condition *test* to the metavariable *?cond*, the context *kont* to *?kont* and the stack *lkont* to *?lkont*. We further restrict condition *?cond* to contain the 'callee' property, of which we take the name and match it to the 'isCMSAdmin' literal. Next, we skip some states until we find a call to the `addNotice` method. Since we only want to detect these calls within the false-branch, we end the policy with an `endIf` predicate with matching stack and context metavariables.

Another option is to find all calls to the `addNotice` method that follow a return of `isCMSAdmin`. Since we only know that the return value of `isCMSAdmin` returns a value of  $\{Bool\}$ , we are unable to rule out any of the branching options. This will result in false-positives. Listing 7.2 gives an implementation of the policy. We again match the stack and context to metavariables *?lkont* and *?kont*, but this time to indicate the start of a function application. Next we specify that we want to find all return statements within that function application. This is done by indicating that these return statements must follow a node which is not the end of the function application, parametrized with the same metavariables for stack and context. Finally, some states can be skipped before finding a function call to `addNotice`.

---

```

1 //First solution
2 G.skipZeroOrMore()
3 .beginIfFalse({test: '?cond', kont: '?kont', lkont: '?lkont',
4                 properties:{
5                     'isCMSAdmin' : '?cond.callee.name'
6                 }}})
7 .skipZeroOrMore()
8 .fCall({name: 'addNotice'})
9 .skipZeroOrMore()
10 .endIf({kont: '?kont', lkont: '?lkont'})
11
12 //Second solution
13 G.skipZeroOrMore()
14 .beginApply({name: 'isCMSAdmin', kont: '?kont', lkont: '?lkont'})
15 .not()
16     .endApply({kont: '?kont', lkont: '?lkont'})
17     .star()
18 .returnStatement()
19 .skipZeroOrMore()
20 .fCall({name: 'addNotice'})

```

---

Listing 7.2: Policy 1 in JS-QL

**7.2.2 A database is opened only after the master password is checked or when creating a new database**

**7.2.3 Public outputs do not depend on a users's password, unless it has been cryptographically hashed**

**7.3 The Conscript language**

**7.4 Evaluation**

# **Chapter 8**

## **Conclusion and future work**

### **8.1 Summary**

### **8.2 Future work**

# Bibliography

- [1] Jamis Buck. *Writing Domain Specific Languages*. 2006. URL: <http://weblog.jamisbuck.org/2006/4/20/writing-domain-specific-languages.html> (visited on 04/16/2016).
- [2] H. Conrad Cunningham. “A little language for surveys: constructing an internal DSL in Ruby”. In: *Proceedings of the 46th Annual Southeast Regional Conference on XX*. ACM-SE 46. Auburn, Alabama: ACM, 2008, pp. 282–287. ISBN: 978-1-60558-105-7. DOI: [10.1145/1593105.1593181](https://doi.org/10.1145/1593105.1593181). URL: <http://dx.doi.org/10.1145/1593105.1593181>.
- [3] Martin Fowler. *Domain Specific Languages*. 1st. Addison-Wesley Professional, 2010. ISBN: 0321712943, 9780321712943.
- [4] Jim Freeze. *Creating DSLs with Ruby*. 2006. URL: [http://www.artima.com/rubycs/articles/ruby\\_as\\_dsl2.html](http://www.artima.com/rubycs/articles/ruby_as_dsl2.html) (visited on 04/16/2016).
- [5] K. A. Hawick. “Fluent Interfaces and Domain-Specific Languages for Graph Generation and Network Analysis Calculations”. In: *Proceedings of the International Conference on Software Engineering (SE’13)*. Innsbruck, Austria: IASTED, Nov. 2013, pp. 752–759.
- [6] Ecma International. *ECMAScript 2015 Language Specification*. 2015. URL: <http://www.ecma-international.org/ecma-262/6.0/index.html> (visited on 04/16/2016).
- [7] Andrew Johnson et al. *Exploring and Enforcing Application Security Guarantees via Program Dependence Graphs*. Tech. rep. TR-04-14. Harvard University, 2014.